

Turan Taghiev

Collected field notes and analysis on trust & safety, platform governance, and AI governance.

17 PIECES · 2026 · 39 PAGES

<https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

These pieces are written from direct operational experience moderating multilingual game communities. Company and product names are withheld under a confidentiality agreement; operational figures are given in rounded bands rather than exact counts.

CONTENTS

FIELD NOTES

- When to Stop Trusting a System You Built 3
- How Long Should Evidence Last: What Old Pattern Data Is Actually For 5
- Onboarding a Moderator: What You Can Teach and What You Can Only Wait Out 7
- What Moderation Data Is Actually Telling You Before Anyone Else Notices 9
- AI Governance Is Architecture, Not Policy 11
- Screening the Screeners: What Actually Predicts a Good Moderator Hire 13
- Same Model, Different Flags 15
- Nine Languages, One Standard: What a Macro Library Actually Buys You 17
- Reviewing the Reviewers: What a Ban Request Looks Like from the Other Side 19
- Repeat Offenders: The Violation Is the Pattern, Not the Message 21
- Explainable Moderation: The Same Explanation Doesn't Work the Same Way Twice 24
- Governance Reporting from the Frontline: What a Report Looks Like When No One Acts On It 26
- Human-in-the-Loop: What It Actually Looks Like from a Moderation Queue 29
- Multilingual Moderation: Same Rule, Different Risk 32

ANALYSIS

- A2/AD and the South China Sea 34
- The TikTok Ban Is Economic, Not Security 36

ESSAY

- Whose Rules? AI Governance as a Power Struggle 38

When to Stop Trusting a System You Built

No system delivers infinite efficiency, and stopping one you built and are responsible for isn't a decision to make lightly. On what "stopping" actually means when a system only flags rather than acts.

The check that actually works isn't trusting your own read on whether it's still net positive — it's the same principle the system was built around: don't let the person closest to the decision be the only one making it.

No system delivers infinite efficiency, and it was never supposed to. The point of the behavioral risk framework I built wasn't to replace judgment, it was to make human judgment faster to apply, to save time that would otherwise go into reading every message from scratch. Even when the system ran on its own, generating flags automatically, the final call always stayed with a person. That single design choice, keeping the decision human no matter how automated the detection got, was what gave anyone real control over what the system actually did.

That distinction matters more than it sounds like it should, because it changes what "stopping" a system even means. A system that only flags, and never acts on its own, doesn't need an emergency kill switch in the way a fully automated one would. The stakes of it running with a bug or a bad pattern are lower, because nothing happens without a human reading the case first. What you're actually deciding, when you ask whether to keep using something like this, isn't "is this dangerous right now," it's "is this still worth the maintenance and oversight it requires."

That's a different question, and it has a different bar. Stopping something you built and are responsible for isn't a decision you make lightly, and it shouldn't be. If a system is doing more good than harm, overall, across the volume of cases it's actually helping with, the imperfections don't automatically disqualify it. No detection system is going to be flawless. Some false positives, some missed cases, some patterns it wasn't built to catch, that's the baseline cost of using pattern-based detection at all, not a sign that the whole thing has failed.

The harder part isn't building that tolerance for imperfection. It's being honest with yourself about where the line actually sits, because you're evaluating something you made, and that comes with its own kind of blind spot. I built this system. That fact doesn't make me more objective about when it's failing, if anything it makes it harder, because every flaw I find is also a flaw in something I designed. The check that actually works isn't trusting your own read on whether it's still net positive. It's the same principle the system itself was built around: don't let the person closest to the decision

be the only one making it. A system that flags instead of acting, and routes every real call to a human, isn't just a safeguard for the players it's protecting. It's a safeguard against the builder's own certainty that everything is still working the way they think it is.

Suggested citation: Taghiev, T. "When to Stop Trusting a System You Built." Field Note, August 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

How Long Should Evidence Last: What Old Pattern Data Is Actually For

How long you keep evidence on a user isn't a fixed number: it depends on what the case file is actually for. On the difference between an open investigation and evidence from a closed case.

Evidence retention isn't about how long you're allowed to hold something against someone — it's about not losing the one thing that turns a scattered set of incidents into a visible pattern.

How long you keep evidence on a user isn't a fixed number. It depends on what the case file is actually for, and that's a question worth asking explicitly instead of defaulting to either extreme, keep everything forever or wipe it after every closed case.

An open, active investigation is the simple version. If a file is being built from a specific starting point, tracking an ongoing pattern, it stays live until that process actually concludes, because it hasn't been used for anything yet and the case is still current. That's not really a retention question. It's just an unfinished file.

The harder version is what happens to evidence after a case closes, after it's already been used for a warning or an action. The instinct might be that it's spent, used once, filed away, done. That's true for one purpose and not for another. The same evidence shouldn't be recycled to justify a second, separate action, that would be double-counting a single incident as though it happened twice. But if the same user surfaces again later, with a new pattern that looks like the old one, the closed case doesn't stay irrelevant. It becomes context. Not standalone proof of anything new, but a note that says this isn't the first time this shape of behavior has shown up, which is exactly the information you need to tell a coincidence from a pattern.

There's a specific reason this matters more than it might seem to. Behavior doesn't actually change on its own. It gets worn down, gradually, through consistent consequences, but the underlying tendency doesn't disappear just because someone stopped acting on it for a while. What can change, though, is the rules themselves. When a policy gets updated, a threshold shifts, a category gets redefined, behavior that used to trigger a clear response can suddenly fall into a gap the new rules don't quite cover. That's not the person changing. That's the framework changing around them, and it can hand back room to maneuver that consistent enforcement had closed off.

That's exactly why old pattern evidence doesn't get discarded just because the rules moved. It's not there to punish someone twice for the same thing. It's there so that if the same shape of behavior resurfaces, under whatever the current rules happen to be, there's a documented reason to look closer instead of treating it as a first offense under a new framework. Evidence retention, done properly, isn't about how long you're allowed to hold something against someone. It's about not losing the one thing that turns a scattered set of incidents into a visible pattern, especially once the rules have shifted enough that the pattern might otherwise look like it's starting over.

Suggested citation: Taghiev, T. "How Long Should Evidence Last: What Old Pattern Data Is Actually For." Field Note, August 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Onboarding a Moderator: What You Can Teach and What You Can Only Wait Out

New moderators start with the same document: rules, tools, escalation paths. That part trains predictably. What doesn't is the instinct not to argue with an angry user, and the deeper judgment about when to hold authority back.

What you can actually do isn't force the slow part to move faster — it's build a structure that gives it something to develop against, instead of sitting there waiting for time to pass.

New moderators get the same starting material: a document covering the rules, the tools, the escalation path. That part is straightforward, and it produces a predictable result. In the first weeks, someone with no prior moderation experience tends to be cautious and rule-following almost by default, because they don't yet have the confidence to deviate from what's written down. That's a fine place to start.

The mistake that shows up most often isn't a misapplication of a rule. It's something the rulebook doesn't cover at all: a new moderator getting pulled into defending the game or the company when a user is venting. It looks like the right instinct in the moment. It's not. Once a moderator starts arguing back, they're no longer managing the server, they're a participant in the conflict, and that changes what they're capable of doing. The exchange escalates, and now there's a second layer of risk: a moderator under pressure, mid-argument, is far more likely to make a real enforcement mistake than one who stayed outside the conflict entirely. The job was never to win the disagreement. It was to keep the space functional. That's a hard thing to explain in advance because it requires believing it before you've felt the pull to respond.

Some of what makes a moderator ready doesn't train on a fixed timeline at all, and that's the harder thing to plan around. The people who actually want this role tend to be high-energy, engaged, opinionated, which is exactly what makes them notice problems in the first place, and exactly what makes them more likely to overstep while they're still calibrating. Given how limited their formal authority is early on, a small misjudgment can carry disproportionate consequences. Some people settle into good judgment in weeks. Others take considerably longer, and there's no reliable way to know in advance which case you're looking at.

What I found actually worked wasn't trying to compress that timeline. It was changing the structure around it. Bringing in two candidates for the same track at the same time, instead of one, does something a training document can't: it turns an open-ended,

unmeasurable development process into something with a visible reference point. Neither person is just improving in the abstract anymore. They're improving relative to someone else going through the identical thing, and that comparison tends to accelerate exactly the kind of judgment that can't be taught directly, because there's now a concrete signal for what "further along" actually looks like.

The rules and the tools can be taught in a document. The instinct not to argue with an angry user takes longer, and mostly needs to be caught and corrected in the moment rather than explained in advance. And the deeper judgment, knowing when authority is real versus when it's better held back, doesn't train on a schedule at all. What you can actually do as the person running onboarding isn't force any of that to move faster. It's build a structure, like parallel candidates, that gives the slow part something to develop against instead of just sitting there waiting for time to pass.

Suggested citation: Taghiev, T. "Onboarding a Moderator: What You Can Teach and What You Can Only Wait Out." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

What Moderation Data Is Actually Telling You Before Anyone Else Notices

When evidence is incomplete, the instinct is to treat the case as unusable. But deleted messages don't disappear from the record entirely: they become a documented possibility rather than proof.

The slow part of a case isn't indecision — it's the difference between having a documented incident and having an established pattern, and only the second is something you can act on with real confidence.

The instinct when evidence is incomplete is to treat the case as unusable. A message gets deleted, a conversation gets edited, and the natural conclusion is that whatever happened there is now unprovable. That's not quite right, and the distinction matters for how you build a case over time instead of only in the moment.

Deleted evidence doesn't disappear from the record entirely. It becomes something with a lower status, not proof, but a documented possibility, a note that something matching a specific pattern happened here, even if the exact content can't be reproduced anymore. On its own, that note proves nothing and shouldn't be treated as if it does. But it doesn't vanish either. If the same pattern shows up again later, in a context where the evidence is intact, the earlier note becomes relevant again, not as standalone proof, but as context that turns an isolated incident into a visible pattern.

That's the real value of holding onto incomplete evidence: not to build a case out of gaps, but to recognize a recurrence when it happens, instead of treating every incident as if it's the first time.

A single message almost never tells you enough on its own, for a related but different reason. Without the surrounding conversation, one line can read as far more or less serious than it actually was. A harsh message in response to genuine provocation reads differently than the same words sent unprompted. Building a case around one screenshot risks getting the severity wrong in either direction. What actually holds up is multiple pieces, from different moments, ideally from different people, that point in the same direction. Not because more evidence is inherently better, but because a pattern across independent sources is much harder to explain away as a misread single moment.

There's also a quieter dynamic that shapes how cases actually develop over time. Someone behaving badly inside a team rarely keeps that fully contained. It tends to

surface, in complaints to peers, in how they talk about the situation to people outside the immediate circle, in the kind of resentment that makes other people willing to come forward. That's not a strategy anyone has to engineer. It's closer to an observation: sustained bad conduct creates friction, and friction generates its own witnesses over time, independent of any single investigator going looking for them. A pattern that seemed hard to prove in month one often has considerably more supporting evidence by month three, not because anyone worked harder to find it, but because the behavior itself kept generating new instances and new people willing to describe what they saw.

That's also why building a real case is slow in a way that acting on it isn't. Establishing a pattern solidly enough to be defensible can take weeks or months of accumulation, formal warnings, documented incidents, corroborating accounts arriving on their own timeline. Once that threshold is actually met, the decision itself, and any procedural action that follows from it, tends to move fast, sometimes within days. The slow part isn't indecision. It's the difference between having a documented incident and having an established pattern, and only the second one is something you can act on with real confidence.

None of this is about lowering the bar for what counts as evidence. It's about not discarding the parts of a case that don't meet that bar yet, because a case rarely arrives complete in a single moment. It accumulates, and knowing what to hold onto while it does is a different skill than knowing what to act on once it's there.

Suggested citation: Taghiev, T. "What Moderation Data Is Actually Telling You Before Anyone Else Notices." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

AI Governance Is Architecture, Not Policy

A governance document can say all the right things and still not govern anything, if none of it is enforced by how the system is actually built. On watching a behavioral risk framework move from written principle to structural rule.

A policy that says humans decide is a sentence; a system where a flag cannot become an action without a logged human decision attached to it is architecture — and the difference only shows when everyone is tired, busy, or under pressure to move fast.

A governance document can say all the right things and still not govern anything, if none of it is actually built into how the system works. I've seen this gap close from the inside, watching a behavioral risk framework move from a written set of principles to something that actually controlled what happened when a flag fired.

The principle everyone writes down is some version of “a human makes the final decision.” That sentence is easy to put in a document. What's harder, and what actually matters, is whether the system is built so that sentence is true by construction, not just true on paper.

The first place this shows up is decision authority, and it isn't flat. Not every flag goes to the same person. Lower-confidence signals route to a regular moderator, with a review window measured in days. Stronger, combined signals route to a senior moderator, on a much shorter clock. The highest-severity patterns, the ones with real potential for harm, require both a senior moderator and a team lead before anything happens, and that requirement isn't a suggestion in a handbook, it's how the case routing is built. Who is allowed to act is tiered to how much damage being wrong would do, and that tiering is enforced by the system's structure, not by moderators remembering to escalate appropriately.

The second is logging, and the detail that actually matters is what counts as a decision at all. Every flag, every review, every outcome gets a timestamp, a reviewer ID, and a stated reason. An action without a log entry isn't treated as an unusual case. It's treated as invalid, as if it didn't happen. That's a stronger claim than “we keep records for audit purposes.” It means the log isn't a byproduct of the process, it's part of what makes an action count as a legitimate one in the first place.

The third is scope, stated as a hard boundary rather than a guideline: the framework covers in-game communication and nothing else. Any data collected for it can't be repurposed for anything outside moderation and trust-and-safety review without separate, explicit approval. This is a constraint built into what the system is allowed to

look at and use, not a footnote about privacy best practices.

There's also a rule that risk signals can't be built from language, nationality, religion, or community identity alone; the signal has to come from behavior. That's not a value statement sitting next to the technical spec. It's a design constraint on what's allowed to function as an input in the first place, which is a different and much stronger thing than a principle a team is expected to remember.

None of this framework was designed in a single sitting from a blank page. Parts of it were formalized after something happened that the existing structure had no answer for, a flood of activity from a single source with no protocol built for that volume. What had been one person's improvised response under pressure became a named, repeatable procedure specifically so the next time didn't depend on the right person being on shift. That's governance built the way architecture gets built: not fully specified in advance, but each real gap closed by turning what worked into a structural rule instead of leaving it as institutional memory.

What separates a governance document from governance itself is whether any of this is actually enforced by the system, or whether it's a description of what people are supposed to remember to do. A policy that says humans decide is a sentence. A system where a flag literally cannot become an action without a logged human decision attached to it is architecture. The difference isn't philosophical. It's whether the rule holds when everyone is tired, busy, or under pressure to move fast, which is exactly when unenforced policy quietly stops being followed and enforced architecture doesn't.

Suggested citation: Taghiev, T. "AI Governance Is Architecture, Not Policy." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Screening the Screeners: What Actually Predicts a Good Moderator Hire

When evaluating candidates for moderation roles, the instinct is to start with what a candidate says about their own judgment. This piece takes the opposite order: verified data first, secondhand accounts only to fill real gaps.

Past behavior is not reset by a new title — it isn't proof of what someone will do, but treating it as irrelevant at the moment of promotion is one of the more common mistakes in deciding who is ready for authority.

When I got asked to weigh in on moderator candidates, the instinct most people have is to start with what the candidate says about themselves, in an interview, in a CV, in how they describe their own judgment. I learned to do the opposite. What someone says about their own conduct is secondary. What I actually look at first is data I can verify myself, and only when there's a real gap or an open question do I turn to what other people say about the candidate.

The clearest version of this shows up with behavioral history. Some candidates have a documented record from before they held any position of authority, complaints, warnings, patterns of how they treated other users when they had no formal role at all. The common assumption is that a title changes a person's behavior, that giving someone responsibility makes them rise to it. Sometimes it does. But a documented pattern doesn't expire just because someone's status changes. If a person was dismissive, hostile, or inconsistent as a regular user, that's not erased by becoming staff. It's information that stays live, and in more than one case, it turned out to predict exactly how that person behaved once they actually had authority to use.

That's the core signal worth naming clearly: past behavior is not reset by a new title. It's not proof of what someone will do, but it's not nothing either, and treating it as irrelevant the moment someone gets promoted is one of the more common mistakes in evaluating who's ready for a moderation role.

The harder skill isn't spotting the obvious red flag. It's knowing when to stop trusting your own read and go looking for a second source. If a candidate's documented history is thin, or a specific incident is ambiguous on its own, that's exactly the point where other people's observations become useful, not as the primary evidence, but as a way to fill a real gap rather than confirm a conclusion I'd already reached. The ordering matters. Verified data first, secondhand accounts only where there's an actual blank to fill, not as a shortcut to skip building the fuller picture myself.

None of this made me the decision-maker in these cases. My role was to observe, document, and flag what I was seeing, and pass that up to whoever actually held the authority to act on it. That distinction matters more than it sounds like it should: the value of careful observation doesn't depend on being the one who gets to decide. It depends on the observation being accurate enough that whoever does decide is working from something real.

Suggested citation: Taghiev, T. "Screening the Screeners: What Actually Predicts a Good Moderator Hire." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Same Model, Different Flags

A moderation filter tuned for English and Russian was applied by default to a Turkish community it had never been calibrated for. The clearest failure case was nationalism: one pattern of speech, two readings, and a filter taught only one of them.

A system built for two languages doesn't degrade gracefully into a third — it applies the wrong pattern with full confidence, and the only way anyone finds out is if someone fluent in what the system can't see is actually looking.

The moderation system I worked with was tuned for two languages: English and Russian. Everyone else's traffic ran through those same settings by default, including a Turkish-speaking community large enough to need its own judgment, using a filter that had never been calibrated for it.

The clearest case of what that actually produces was nationalism. In the global, English-language channel, nationalist language reads as close to racism, tribal, exclusionary, often adjacent to language the filter was already tuned to catch. In the Turkish channel, similar language regularly read as patriotism, something other Turkish users would recognize, agree with, and rally around rather than report. Same underlying pattern of speech. Two communities, two completely different readings of what it meant, and one filter that had only ever been taught one of those readings.

This isn't a claim that the model was "wrong" in either direction. A filter tuned on English-language moderation data learns English-language patterns of what nationalist language tends to escalate into. It has no equivalent training signal for how the same category of speech functions in a Turkish community, because nobody built one. The gap wasn't a bug in the model. It was a scope the model was never given.

What that produced in practice was inconsistent, not because the system was unreliable, but because it was being asked a question outside what it had been calibrated to answer. Language that would have been flagged without hesitation in English sometimes passed through untouched in Turkish, not because it was less risky, but because nothing in the filter's training had told it to look for that pattern in that language. The reverse happened too: phrasing that read as ordinary, even mild, competitive banter in a Turkish gaming context occasionally tripped thresholds calibrated on English norms, because the filter was pattern-matching against the only language model it actually understood well.

The fix wasn't adjusting the existing filter's sensitivity. Turning it up or down wouldn't have helped, because the problem wasn't calibration within the language it knew. It was the absence of a language it didn't. What actually closed the gap was building a

Turkish-specific layer from the ground up, a blacklist and pattern set constructed around how nationalism, trash talk, and genuine harassment actually show up in that specific community, informed by someone who could tell the three apart in context the filter couldn't parse. Every case that layer caught, and every case it correctly let through, ran through manual review before it went anywhere close to full trust in the automated flag, precisely because the new layer hadn't earned the calibration history the English/Russian setup had.

The broader point isn't specific to gaming or to this one filter. Any AI moderation system inherits the linguistic and cultural scope of whatever it was trained and tuned on, and that scope doesn't expand automatically just because the platform's user base does. A system built for two languages doesn't degrade gracefully into a third. It just applies the wrong pattern with full confidence, and the only way anyone finds out is if someone fluent in what the system can't see is actually looking.

Suggested citation: Taghiev, T. "Same Model, Different Flags." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Nine Languages, One Standard: What a Macro Library Actually Buys You

Players in uncovered regions kept asking why support didn't exist in their language. The staffing fix wasn't available on that timeline, so what got built instead was a macro library across seven additional languages, run solo for over a year.

A macro doesn't remove judgment from the interaction — it removes the part that never needed judgment in the first place, so what's left can go where it's actually required.

The complaint that started this wasn't about moderation quality. It was simpler: players from certain regions kept asking why there was no support in their language, while players in a handful of major languages got answered quickly. The gap wasn't in the rules. It was in coverage. We had moderators fluent in exactly two languages beyond English, and a global player base that didn't map onto that at all.

The obvious fix, hiring native-language helpers for every region, wasn't available on the timeline the complaints were happening on. So I built the thing that was available: a macro library covering seven additional languages, ten response templates each, translated and quality-checked, so that a moderator who didn't speak a word of Romanian or Portuguese could still respond to a player in that language without the response reading like it came from a translation tool. I paid for the AI translation tooling myself, out of pocket, because the alternative was leaving those complaints unanswered until a hiring process caught up. I ran that system alone for over a year.

The obvious worry with any macro system is that it makes everything feel robotic; the same ten answers, over and over, in every language, until players notice they're talking to a script. That worry turned out to be about the wrong layer. The response text was standardized. The decision about which response to use, when to deviate from it, when a case needed something the macro didn't cover at all, that stayed human, because I was the one making it every time. Most tickets fell into two or three recognizable categories: the same handful of questions that spike after every content update, phrased slightly differently by every player who asks them. A macro doesn't remove judgment from that interaction. It removes the part of the interaction that never needed judgment in the first place, so the judgment that's left can go toward the smaller number of cases that actually require it.

The distinction became clearer once I built the support-line version of the same idea, a separate set of macros specifically for technical support tickets, built off actual complaint patterns rather than a general translation pass. Support tickets cluster hard around whatever broke in the last update: the same three or four issues, arriving in bursts, from players who have no way of knowing thousands of other people are asking the identical question that hour. Answering each one individually, from scratch, at that volume, isn't more thoughtful. It's slower for the player and it doesn't produce a better answer. What changes with volume isn't how much judgment a response needs. It's how much of that judgment has already been made correctly the last time, and doesn't need to be re-derived from nothing every single time it comes up again.

What the system actually bought wasn't speed for its own sake, and it wasn't uniformity for its own sake either. It was the ability to say, to players in regions the team wasn't staffed to cover: support exists here too, and it isn't going to feel like an afterthought just because there's no regional helper yet. Coverage and consistency turned out to be the same problem wearing two names. Once every language had a floor it could rely on, the actual variation in how each case got handled came entirely from what a human decided to do with it, not from which language the player happened to type in.

Suggested citation: Taghiev, T. "Nine Languages, One Standard: What a Macro Library Actually Buys You." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Reviewing the Reviewers: What a Ban Request Looks Like from the Other Side

Most moderation writing is about deciding what to do with a user. This piece is about the layer above: reviewing another moderator's ban request before it becomes final, against three criteria none of which were about the ban itself.

A ban request that fails review isn't necessarily wrong; it just isn't ready — sending it back corrects the case the helper will need to defend, not their judgment about the user.

Most writing about moderation is written from the perspective of the moderator deciding what to do about a user. Less gets written about the layer above that: reviewing someone else's decision about a user, before it becomes final.

I spent time reviewing ban requests submitted by a volunteer helper team, moderators one step below full authority, who could warn, mute, and escalate, but not issue permanent bans on their own. Every request that reached that ceiling came through me first. What I was actually checking wasn't just whether the user deserved the penalty. It was whether the case in front of me could survive being questioned months later.

The standard I checked against had three parts, and none of them were about the ban itself.

The first was evidence, not conclusions. A helper could be completely right about a user and still submit a request I'd send back, because "this user was clearly trading accounts" isn't evidence, it's a summary. What I needed was the actual message, the actual timestamp, the actual pattern across multiple incidents if there was one. The reason wasn't procedural box-checking. It was that a banned user has months to come back and claim they were punished for nothing, and if the evidence behind that ban is thin or missing, the action gets reversed by whoever reviews it next, which damages the helper's credibility more than not banning the user at all would have.

The second was whether the escalation ladder had actually been climbed, not skipped. The path was supposed to run from warning to formal warning to mute to ban, and I was checking whether a helper had jumped straight to the top because a case felt serious in the moment, versus whether the severity genuinely justified skipping steps. Both happen. Only one of them is defensible when someone asks why.

The third was the part that took the longest to learn to recognize, and it wasn't really about rule violations at all: how a helper had handled a user who was operating right at

the edge of the rules, provocative, borderline, clearly testing limits, but not yet clearly over the line. The instinct for a newer helper is to act early, because the behavior is obviously heading somewhere bad. The problem is that acting early on a borderline case, before there's a violation a reasonable person would call unambiguous, hands that user exactly what they need. A moderator who penalizes something that wasn't quite a violation yet becomes the story, not the user's behavior. I'd send those requests back not because the helper's read on the user was wrong, but because the timing would have made the helper the problem instead of the solution. The better move, and the harder one, was watching, documenting, and waiting for a violation nobody could argue with, then responding to that one clearly and, often, more firmly than a first offense would normally get, because the pattern leading up to it was now part of the record too.

None of this was written down as a formal rubric. It became one anyway, because reviewing enough of these requests makes the actual criteria visible even when nobody articulated them in advance: is the evidence specific enough to survive a challenge, was the escalation path respected or skipped, and did the timing protect the helper or expose them. A ban request that failed any of the three wasn't necessarily wrong. It just wasn't ready, and sending it back wasn't a correction of the helper's judgment about the user. It was a correction of the case they'd need to defend it.

Suggested citation: Taghiev, T. "Reviewing the Reviewers: What a Ban Request Looks Like from the Other Side." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Repeat Offenders: The Violation Is the Pattern, Not the Message

A month of hand-logged moderation data showed most violators appear once and disappear, while a small group returns again and again. Per-incident moderation handles every case correctly and still gets the person wrong.

The unit of moderation is the incident, but the unit of risk is the person — and anywhere that gap is bridged only by a moderator's memory, the community is one staffing change away from losing its pattern detection.

Most moderation writing treats a violation as an event: a message appears, it breaks a rule, someone acts on it, the case closes. The queue is built around that unit. So are the logs, the reports, and most of the tooling. After a month of moderating a Brazilian Portuguese game community and logging every violation by hand, I think that unit is wrong, or at least incomplete. The message is what you act on. The person is where the risk lives.

Here is what a month actually looked like. Across thirty days and about 5,700 messages, I logged roughly 130 violation cases. They came from about 70 unique users. That ratio already tells you something: most people who break a rule do it once, get warned or muted, and never show up in the log again. But roughly one in five violators came back at least once. And a much smaller group came back again and again, across different days, with the same behavior each time.

One case made the problem impossible to ignore. A user asked to buy an account in the channel. Account trading was a clear violation, so the message was removed and the user warned. The next day, the same request, twice. The day after, again. By the time the same user posted the same request on a fifth separate day, the individual messages had not changed at all: each one, taken alone, was a mild, almost polite violation, the kind that earns a routine warning. What had changed was everything around them. This was no longer a user who broke a rule. It was a user whose relationship to the rule was the violation.

Per-incident moderation handles each of those days correctly and still gets the month wrong. Every removal was justified. Every warning followed procedure. If you audited any single case, you would find nothing to criticize. But the correct handling of five incidents added up to an incorrect handling of one person, because the escalation logic lived in my memory, not in the system. I recognized the name. I remembered the previous days. The ban recommendation that eventually went forward was built on continuity that existed nowhere except in a moderator who happened to have been on shift all week.

That is the structural weakness. A moderation system that tracks incidents but not people delegates its pattern detection to whoever happens to be watching. It works when the same moderator covers the same channel every day, holds the names in their head, and writes careful logs. It fails silently the moment coverage rotates, volume spikes, or the moderator who carries the informal database leaves. The pattern does not disappear. The ability to see it does.

The repeat cases also behaved differently from the one-off cases, in a way that matters for enforcement design. First offenses in that month were scattered across categories: spam, self-promotion, off-topic begging, an occasional heated exchange. The repeat offenses concentrated. They clustered heavily around account trading and adjacent commerce: buying, selling, matching buyers with sellers. At one point a second user started tagging the repeat offender to connect a purchase. That is not two isolated violations. That is a small market forming, and a market is a pattern by definition. You cannot see it at the message level, because each message is just one side of one conversation.

This changed how I wrote recommendations. A first-time violator got the action the incident deserved. A repeat name got the action the trajectory deserved, and I started saying so explicitly in the log: not “user posted a prohibited request” but “user posted the same prohibited request for the fifth day; warnings have demonstrably not changed the behavior; recommend removal.” The evidence for a ban was never the fifth message. It was the failure of four previous interventions, documented in sequence.

There is an obvious objection: giving weight to history risks punishing people for who they are rather than what they did. It is a fair concern, and it is why the sequence matters. Escalation earned by documented, repeated, same-category violations is not profiling. It is the opposite: it is the system giving someone four chances and recording each one. What would actually be unfair is the alternative that per-incident moderation quietly produces, where a persistent bad actor and a one-time offender receive identical treatment forever because the system cannot tell them apart.

What this points to. The unit of moderation is the incident, but the unit of risk is the person, and most community tooling only knows about the first. Anywhere the gap between those two is bridged by an individual moderator’s memory, the community is one staffing change away from losing its pattern detection entirely. If I could push one requirement into any moderation stack, it would not be a better filter. It would be a persistent, queryable answer to the oldest question in the queue: have we seen this person before, and for what?

Suggested citation: Taghiev, T. “Repeat Offenders: The Violation Is the Pattern, Not the Message.” Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at

<https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Explainable Moderation: The Same Explanation Doesn't Work the Same Way Twice

Most writing about explainable moderation focuses on how much detail a rule violation notice should include. Less attention goes to tone, and to the fact that tone can't be standardized across communities without cost. Drawing on moderating the same game across Turkish, English, and Brazilian Portuguese communities, this piece treats tone as a resource allocation decision rather than a stylistic one.

Tone is a resource-allocation decision, not a stylistic one — spending it uniformly across communities with different escalation patterns isn't neutral.

Most writing about explainable moderation asks the same question: how much detail should a rule violation notice include. What gets less attention is that the tone of that explanation matters as much as its content, and that tone isn't something you can standardize across communities without cost.

I moderated Turkish, English, and Brazilian Portuguese communities for the same game, often issuing the same category of penalty across all three in a single shift. The written policy was identical. The delivery wasn't, and it wasn't supposed to be.

With Turkish users, I defaulted to neutral, low-affect language, deliberately. Almost any enforcement action, regardless of how clearly justified, would draw an appeal. Not a quiet one. Users would argue they'd been provoked, claim they didn't know the rule, or resurface the same complaint months later insisting the ban was unfair and they'd "done nothing." Some of that pushback stayed civil. Some of it escalated into insults, and a smaller number into outright fabrication, users telling other servers a moderator had a personal vendetta against them, reframing a justified penalty as targeted harassment. Warmth in that environment wasn't a neutral choice. It read as an opening. So the tone I used wasn't really about the individual case. It was risk management: a flat, procedural explanation gave less material to escalate, and conserved energy I'd need for the cases that actually required it.

With English-speaking users, the same neutral tone would have been a mismatch. That community ran calmer by default, disagreement without the personal-vendetta narrative, and enforcement was more often met with quiet acceptance or an actual apology. Being warmer here wasn't a risk. It was accurate to the environment, and it also did something specific: as my role became more visible, that same warmth read differently than it would have in the Turkish community. Authority read as approachable

rather than as an opportunity to escalate.

The Brazilian Portuguese community is the clearest test of whether this was really about language or about something else, because I never spoke the language well enough to modulate tone through word choice alone. Every explanation went through translation. What I could control was intent behind the interaction, staying warm and operationally friendly regardless of the specific words landing correctly. I have not had a single user push back on an enforcement action in that community. That's not because the moderation was lighter. It's because the posture behind the explanation, even filtered through translation, seemed to matter more than its precision.

None of this means one tone is objectively correct and the others are overcautious. It means "explainable moderation" as a design principle assumes there's a single right way to explain a decision, when what I was actually managing across three communities was three different appeal cultures, three different baseline relationships to authority, and three different costs of getting the tone wrong. A flat explanation isn't inherently better or worse than a warm one. It's a resource allocation decision, made under the assumption that moderator attention and goodwill are finite, and that spending them uniformly across communities with radically different escalation patterns isn't neutral. It's just spending them in the wrong place.

Suggested citation: Taghiev, T. "Explainable Moderation: The Same Explanation Doesn't Work the Same Way Twice." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>


Governance Reporting from the Frontline: What a Report Looks Like When No One Acts On It

Most writing about governance reporting focuses on what a good report contains. Less gets said about what happens when a report is accurate, repeated, and correct, and still doesn't change anything. Drawing on four consecutive months of real moderation reports for a community I moderated without speaking its language, this piece looks at what reporting can actually force, and what it can only document.

A recommendation, repeated correctly with mounting evidence, is not the same thing as a decision — reporting can document a gap, but it can't force the fix.

Most writing about governance reporting focuses on what a good report contains. Less gets said about what happens when a report is good, repeated, and correct, and still doesn't change anything. I spent four consecutive months writing monthly moderation reports for a Brazilian Portuguese game community I moderated without speaking Portuguese. What that produced wasn't just a data problem. It was a real-time case study in what a report can and can't do on its own.

The pattern was visible from month one. Every report followed the same structure: activity data, violation breakdown, operational reality, recommendations. The recommendations were consistent across all four months, almost word for word: hire a Portuguese-speaking moderator, build a blacklist for patterns detectable without translation, establish a coordination protocol for high-activity periods. March called moderator hiring “not a comfort decision, a capacity decision.” June called it “not a comfort decision. It is a capacity decision” again, three months later, because nothing had changed.



Volume spiked in April, then eased back toward where it started (rounded)

The data behind the repetition tells its own story. Message volume roughly doubled between March and April (about 6,300 to about 11,900), then eased back through May (about 8,900) and June (about 5,700), nearly returning to where it started. Traffic came back down. The pattern underneath it didn't. Violations followed the same curve: April's

report logged 49 violations across categories from account trading to unauthorized recruitment; May counted 74, with account trading alone making up a third of all cases. By June, the violation types hadn't diversified so much as organized: account and gold trading, gambling-site referrals, and a pay-for-task scheme built around TikTok Lite referral links, structured enough to look like coordinated outside activity rather than isolated bad behavior.

What made repeat offenders hard to catch wasn't the language barrier alone. It was the absence of a system to remember them. The April report documents a user who committed the same violation six consecutive days, each time processed in isolation as a one-hour mute, because there was no tracking mechanism fast enough to catch the pattern under single-moderator conditions. Standard escalation logic (warning, then temporary mute, then longer mute, then ban) would have ended the behavior on day three. Instead it ran for six, not because the moderation policy was wrong, but because applying it correctly required remembering a case that no system was recording. May's report is explicit about this: nine separate ban evasion cases that month were caught only because one person recognized writing style and message patterns from memory. That's not a detection system. That's a moderator functioning as the detection system, and the report says so directly.

The reports also captured something reports don't usually have room for: who else was absorbing the gap. May's report includes a section crediting an unofficial volunteer who had been directing users to the right channels and explaining rules in two languages, alongside a direct statement that this was "a serious structural problem," not a success story. A community filling in for a missing moderation function isn't evidence the system is working. It's evidence of exactly how much slack the current setup depends on to look stable from the outside.

By June, the reports had started measuring their own failure to produce change. Each month's conclusion table scored the same categories: operational stability (maintained), violation detection (consistent), capacity sustainability (at critical limit), language barrier (not resolved), system design (incomplete). The status columns didn't move for four months. What did move was the language around them, from "close to breaking" in March to "past critical limit" in April. A report can hold a system's status constant in its own summary language while the substance underneath it keeps deteriorating; the report's job was never to fix that, only to make it visible in numbers regularly enough that "we didn't know" stopped being a defensible response.

What this points to. A recommendation, repeated correctly with mounting evidence, is not the same thing as a decision. Reporting can establish that a gap exists, quantify how it's growing, and document what specifically breaks when it's not closed, month after month, on a fixed schedule, in language leadership can act on. What it can't do is force the action itself. That's a different function, with a different owner, and conflating the two is exactly how an operational report turns into an accountability record instead: not proof that the system worked, but proof of exactly when someone knew it wasn't going

to.

Suggested citation: Taghiev, T. "Governance Reporting from the Frontline." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Human-in-the-Loop: What It Actually Looks Like from a Moderation Queue

“Human-in-the-loop” is usually presented as a safeguard: the system flags, a person decides, harm is prevented. Two and a half years inside a live moderation queue for a multiplayer game show a messier picture: context collapses into single keywords, technical false positives punish innocent messages, and review capacity itself becomes a constraint that governance design has to plan for.

Human-in-the-loop is a capacity constraint, not a guarantee — and the least-discussed part of the principle is what happens when review capacity runs out.

Most AI governance writing treats “human-in-the-loop” as a settled answer: the system flags something, a person reviews it, and that review is the safeguard. I spent two and a half years moderating and doing technical support for a multiplayer game’s Turkish and Brazilian communities, working with two versions of an internal chat-moderation system. What I want to describe here isn’t the policy version of human-in-the-loop. It’s the queue version: what the principle actually looks like once it meets real volume, real ambiguity, and a keyboard.

The word problem. The first system we used was a keyword blacklist: five languages, ten violation categories, a severity scale from a mild toxicity flag up to an immediate block. It worked fine for the obvious cases. A slur is a slur, a threat is a threat. The problems showed up at the edges, where a word’s presence didn’t match the actual intent behind it.

One in-game item type could be bought through unofficial, off-platform sellers, a scam players warned each other about. The word for that item sat on the blacklist, because scammers used the same word to lure people in. But a player warning someone not to buy through those channels used the exact same word. The filter couldn’t tell a warning from a pitch. Both got blocked.

That didn’t just produce a wrongly flagged message. It changed how people talked. Players got more careful, more likely to leave a sentence unfinished rather than risk a mute they’d have to appeal. Multiply that across a chat channel and the server gets quieter, not because the rule removed bad actors, but because it made normal players self-censor.

The second kind of false positive was simpler, and honestly more revealing. Any string that looked like a URL got flagged as a link, a reasonable default for scam prevention.

But Turkish punctuation habits don't always leave a space after a period. Someone typing "iyiyim.Sen nasılsın" ("I'm fine. How are you") with no space after the full stop produced something that read like a domain name. The system treated it the same as a phishing link. There was no judgment call here, no context to weigh. Just a filter that didn't account for how a period followed by a lowercase letter could mean something completely different in Turkish. This is where "human-in-the-loop" earns its keep in the most literal way: not making a hard call, just recognizing the machine got something obviously wrong, and having the standing to fix it.

Override authority was never a fixed rule. It depended on seeing the full message in context, not just the isolated flagged word. That gap was exactly what the next version of the system was built to close.

From a single message to a pattern. The system we moved to didn't replace the blacklist. It sat on top of it and asked a different question. Instead of "did this message contain a forbidden word," it asked "does this account's behavior, read across sessions, look like it's heading somewhere." One insult in a heated match is trash talk. The same player escalating from "nice play" to personal attacks against the same target over five straight messages is a different signal, even though no single line in that sequence would trip the original filter. An account sending compliments over several sessions, then proposing a move to Discord, then offering "insider access" is a scam being built one harmless-looking step at a time.

In practice, this meant comparing where someone's behavior seemed to be heading against known risky patterns, and routing the matches to a person before anything happened, not after. That's a genuinely different kind of problem than keyword matching, and it's a more honest use of "AI-assisted" moderation, because the system's output was never an action. It was a routing decision. It decided who a human should look at next, not what should happen to them. That distinction, flag versus verdict, is really the whole point of "human-in-the-loop," and it's worth taking seriously because it's easy to say and easy to quietly lose.

Where the loop actually breaks. The clean version of this story is that every flag gets someone's full attention. The honest version is that review capacity is limited, and any real system has to say what happens when it runs out.

I didn't learn this from a policy document. I learned it on a weekend with no senior staff on shift, when a single account generated somewhere between six hundred and seven hundred support tickets in a short window. Repetitive content, accelerating fast, no existing protocol for it. There was no time to open six hundred tickets and read each one individually. What actually happened was triage: find the account through the support platform's search, close the whole cluster under one resolution, absorb the volume into that shift, and flag the account to look at more closely later. That's still a human making a decision. It's just not the decision the governance language implies when it promises individual review of every case. It's closer to someone deciding, correctly, that

individual review was the wrong tool for this particular problem, and using a rougher one on purpose.

What happened after is the part that matters most. That improvised response got written down. The volume threshold, the pattern, the containment steps all became a named, repeatable procedure, so the next flood wouldn't depend on the right person happening to be on shift that day.

What this actually shows. Not that human review is decorative, and not that it's enough on its own. Both of those are easy, wrong answers. The real point is narrower: "human-in-the-loop" isn't something you get to claim once and stop thinking about. It's a capacity constraint, and it behaves like one. It holds under normal load and bends under a flood, the same as any staffing model would. Systems that promise human review without planning for the weekend six hundred tickets show up at once aren't describing a safeguard. They're describing a hope that happens to use safeguard language.

The version of human-in-the-loop that actually survived contact with the queue wasn't the one where a person reads every flag with equal care. It was the one where someone had the standing to override an obviously wrong call, the tools to see behavior as a pattern instead of one isolated message, and, most importantly, an explicit, already-approved way to triage at scale when individual review wasn't possible. That third part usually gets left out of how this principle is discussed. It shouldn't be. It's the part that gets tested first.

Suggested citation: Taghiev, T. "Human-in-the-Loop: What It Actually Looks Like from a Moderation Queue." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Multilingual Moderation: Same Rule, Different Risk

The hard part of moderating political speech across languages isn't the ruling. It's that the same enforcement lands as fair in one community and as betrayal in another, and repeated correct calls can turn a moderator into a target.

A rule enforced correctly ten separate times can still end a moderator's standing in the community if no one tracks that it was the same ten people, reacting the same way, every time.

The hardest part of moderating political speech across languages isn't deciding what counts as a violation. The real difficulty is this: the same enforcement action can be perceived completely differently depending on which community is watching.

Take an example. When a political comment is posted in a Turkish channel, other Turkish members usually see it as normal, even expected: a shared cultural reference point, not a threat. The same comment, translated and dropped into the English channel, can read as racism or extremism. Because the English channel serves a global audience, there's no shared national context to soften it.

The standard process is identical on paper: verbal or written warning, timeout, then a temporary or permanent ban depending on severity. Applying this in the English channel feels fair: global audience, universal standard, no exceptions. Applying the same process in the Turkish channel often doesn't feel fair to the people living in it. Why would a Turkish moderator ban a Turkish user for defending their own country? The action is policy-correct, but it lands like a betrayal, and that perception problem is separate from whether the enforcement itself was right.

This is where a second layer matters, just as much as the first: managing community perception without compromising the rule. Short, consistent, public explanations, not defensive, not apologetic, keep the enforcement solid while protecting the moderator's standing. But if this cycle continues, something more important starts happening in the background: the same "patriotic" poster, repeated over and over, can generate a wave of complaints against the same moderator. Not because the moderator did anything wrong, but because enforcing an unpopular-but-correct rule repeatedly turns them into a target.

The critical operational point is this: it's not about counting complaints. Complaint volume against a moderator means nothing on its own, it's a signal, not a verdict. What matters is documenting the underlying behavioral pattern: same users, same violation type, same enforcement logic, timestamped. This documentation exists for one purpose,

protecting the moderator from being effectively pushed out by coordinated backlash, even when no single complaint is malicious. This functions as a form of mobbing.

For less experienced or time-constrained moderators, this step is easy to skip, because each individual case looks resolved on its own. But the real risk is the pattern. A rule enforced correctly ten separate times can still end a moderator's standing in the community if no one tracks that it's the same ten people, reacting the same way, every time.

Suggested citation: Taghiev, T. "Multilingual Moderation: Same Rule, Different Risk." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

A2/AD and the South China Sea

China is not trying to win a fleet battle. It is trying to make American sea control too expensive to use.

The balance in the South China Sea has not flipped, but it has shifted at the level of will and continuity rather than raw capability — and the trend line only points one way until the US shows it will pay to reverse it.

The South China Sea carries about 3.4 trillion dollars of trade a year, along with a fifth of the world's oil, and the question of who controls it is being settled right now. The usual way to read the contest is capability against capability: count China's ships, count America's carriers, compare. That framing misses what China is actually doing.

China is not building a navy to defeat the US Navy. It is building an architecture that makes American sea control too expensive to exercise. The DF-21D missile puts carrier groups at risk out to 1,500 kilometres, covering the entire First Island Chain; the DF-26 reaches Guam at 4,000, and carries both conventional and nuclear warheads. The artificial islands in the Spratlys, some 1,300 hectares of reclaimed land since 2013, extend radar coverage and missile range into the middle of the sea. A quiet diesel-electric submarine fleet, hard to detect in the sea's shallow, thermally layered waters, makes every US operational plan start with an anti-submarine problem. Jamming stations and anti-satellite weapons threaten the information layer everything else depends on. None of this is built to win a battle. It is built to change the maths before any shot is fired.

There is a historical name for this. Before the First World War, Admiral Tirpitz built a German fleet that could never beat the Royal Navy but could make British naval dominance too costly to enforce. Sea denial, not sea control. China has taken the same logic and rebuilt it with missiles, islands and electronic warfare instead of battleships. The aim is narrow and precise: not to change what the US Navy is, only to change what it calculates before entering one specific body of water.

The missiles are the visible layer. The daily work happens lower. China's coast guard, now one of the largest and most heavily armed in the world, blocks Philippine resupply runs with water cannon and lasers. A maritime militia of nominally civilian fishing boats keeps a permanent presence in disputed waters. No single incident crosses the threshold that would justify a military response, and any response to a "civilian" vessel carries a legitimacy cost. Each move is small. The accumulation is the strategy.

The American reaction has been real but limited. Freedom of navigation operations keep the legal claim alive without changing anything on the water; Beijing files a protest after each one and continues building. The doctrinal shift from AirSea Battle to

JAM-GC adapted operational planning to a denial environment. The alliance network has thickened: QUAD, AUKUS, expanded base access in the Philippines. Japan is now building its own denial layer along the Ryukyus, turning the logic back on Beijing.

The 2016 Hague ruling declared China's nine-dash-line claims invalid under UNCLOS, and construction continued the next day. There is also a quieter problem underneath the legal argument: the United States builds its case on UNCLOS, a treaty it has never ratified. Beijing does not miss the opening.

The balance in the South China Sea has not flipped. But it has shifted, and the shift happened at the level of will and continuity rather than raw capability. China made its presence permanent, physically and operationally. The US has not yet shown it is willing to pay what reversing that would cost. Until it does, the trend line only points one way.

Suggested citation: Taghiev, T. "A2/AD and the South China Sea." Analysis, 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

The TikTok Ban Is Economic, Not Security

The official explanation is national security. The evidence points somewhere else. Between 2020 and 2024, the ban tracked market share far more closely than any intelligence finding.

The security hypothesis explains TikTok alone, with evidence that stays theoretical; the competition hypothesis explains TikTok, Huawei, and export controls with one logic.

The official story about the TikTok ban runs like this: a Chinese company is vacuuming up American user data, and under China's 2017 National Intelligence Law, that data could end up in Beijing's hands. It's a plausible story. The problem is what happens when you test it.

In 2020, the CIA and the Justice Department looked into the claims. They came back without concrete evidence. The algorithm might carry a risk. The data collection is real. But the same collection happens at Meta and Google, and Washington isn't talking about banning them. If the concern were really about user data, the obvious move would be a federal privacy law covering every platform. The US still doesn't have one.

The strongest evidence against the "one administration's obsession" reading is continuity. Trump demanded a sale or shutdown in 2020. Biden paused parts of that push, then Congress picked it up with new bills in 2023 and 2024, backed by both parties. Administrations changed; the pressure didn't. Policies that survive a change of government usually aren't personal preferences. They're structural interests.

TikTok's own response makes the same point from the other side. Under Project Texas, the company moved American user data to US-based storage with outside oversight, exactly the kind of technical fix a security problem calls for. The pressure didn't ease. When a company addresses the stated concern and the policy doesn't move, the stated concern probably isn't the operative one.

What the timeline does track, quite closely, is market share. Between 2020 and 2024, TikTok grew from a novelty into one of the largest advertising platforms in the country, cutting straight into Meta's territory. Analysts estimated that a ban would add somewhere near \$431 billion in combined market value to Meta, Google, and Amazon. That's not background noise.

The economic motive shows up elsewhere too. ByteDance hit around \$155 billion in annual revenue by 2024, roughly Meta's level. Washington's answer wasn't to regulate digital platforms across the board. It was to force one Chinese company to sell to

American buyers or shut down, with Microsoft and Oracle waiting to buy. That's not how you handle a security threat. It's how you take out a competitor.

The realist reading is the simple one. States move to protect their economic position, and they reach for justifying language when they do. National security supplies that language cleanly. None of this proves the data worries are invented; they're real, and they're still unresolved. It proves that when those worries line up with commercial interests, the commercial interests tend to set the timing and shape of the response.

Put the two explanations side by side and the asymmetry is clear. The security hypothesis explains TikTok, and only TikTok, and only with evidence that stays theoretical. The competition hypothesis explains TikTok, Huawei, ZTE, and the semiconductor export controls with a single logic: limit the rise of Chinese technology capacity wherever it starts to win.

China's closed market looks different from the outside, but the logic underneath is the same. Beijing keeps foreign platforms out and calls it sovereignty. Washington pushes Chinese platforms out and calls it security. Both governments are running the same calculation; they just use different words for it.

The next decade of digital advertising will be worth several trillion dollars. The TikTok fight was, when you strip it down, about who ends up positioned to take it.

Suggested citation: Taghiev, T. "The TikTok Ban Is Economic, Not Security." Analysis, 2025.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>

Whose Rules? AI Governance as a Power Struggle

AI governance is presented as a technical safety problem. It is a contest over who gets to write the rules.

No governance framework can do its job while its legitimacy is contested by one of the two actors it most needs to bind — safety language cannot paper over a hegemony problem.

Every few months a new framework for governing artificial intelligence arrives: an executive order, an EU act, a summit declaration. They differ in detail but share a structure. Western technology is treated as the baseline, the normal thing that needs guardrails. Chinese technology is treated as the threat, the abnormal thing that needs containment. That asymmetry is not an oversight in the design. It is the design.

Read this way, the governance debate stops being a technical conversation about safety and becomes what it actually is: a contest over who writes the rules of the next technological order, and whose systems get defined as legitimate by default.

The TikTok ban is the cleanest evidence. The official justification was national security; the investigations that were supposed to support it came back without concrete findings, while the ban's timeline tracked TikTok's advertising market share far more closely than any intelligence assessment. The companies that stood to gain, Meta first among them, were the loudest voices for restriction. I have written about the economics of that case separately. What matters here is the structural point: a state that preaches open markets closed one, and did it in the language of security.

Beijing noticed, because Beijing runs the same play with different vocabulary. China keeps foreign platforms out and calls it sovereignty. Washington pushes Chinese platforms out and calls it security. From a realist reading both are doing the identical thing, protecting a strategic position, and the fact that only one of them is coded as a violation of liberal norms is precisely the asymmetry Chinese officials point to when they refuse to accept Western governance frameworks as neutral.

Robert Cox drew a distinction that fits this moment exactly. Problem-solving theory takes the existing order as given and works out technical fixes within it. Critical theory asks where the order came from, who it serves, and whether it deserves the authority it claims. Cox's point was that theory is never neutral: it is always for someone and for some purpose.

Western governance proposals debate compute thresholds, evaluation regimes, licensing schemes, all fixes within a framework whose legitimacy is assumed. China's

response operates one level up. It does not offer a better clause; it questions why a framework written by its strategic rival, with its rival's systems positioned as the default, should bind it at all. Whatever one thinks of Beijing's motives, and they are self-interested, the move is analytically coherent. You do not negotiate the terms of an order that defines you as the threat.

There is a historical pattern here that governance optimists keep walking past. Institutions that exclude a major power do not constrain that power; they stop functioning. The League of Nations was built without the United States and could not survive its absence. An AI governance regime built around the containment of China inherits the same defect: the actor with the fastest-growing AI capacity in the world has no reason to comply with rules written against it, and every reason to build a parallel order, which is what it is doing.

The uncomfortable arithmetic follows. A framework that matters requires the United States to give up some of the hegemonic framing, treating Chinese participation as something to negotiate rather than something to prevent, and it requires China to give up some of its closure, offering real market access as the price of a seat. Neither concession is likely soon. Both are structurally necessary. Without them, "global AI governance" will remain what it currently is: two regulatory blocs, each calling its own sphere universal.

None of this is an argument that AI carries no real risks, or that governance is pointless. The risks are real and some of the technical work is serious. The argument is narrower and harder: no governance framework can do its job while its legitimacy is contested by one of the two actors it most needs to bind. Safety language cannot paper over a hegemony problem. Until the power question is addressed as a power question, every summit declaration is a document one half of the world has already decided not to read.

Suggested citation: Taghiev, T. "Whose Rules? AI Governance as a Power Struggle." Essay, 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>