

FIELD NOTE • 2026 • TRUST & SAFETY • AI GOVERNANCE

What AI-Assisted Moderation Actually Caught

The hardest violations to catch are the borderline cases backed by influence, not the obvious ones. No single conversation reveals that on its own.

I never treated its output as the decision. It was an input I checked against the actual case before anything went out.

The hardest violations to catch aren't the obvious ones. They're the borderline cases, where someone knows exactly how close to the line they can get without crossing it clearly, and knows it because they've had backing before, someone willing to argue on their behalf, someone with enough standing to turn a straightforward decision into a fight. That's the category where a longer-term pattern actually matters most, and it's also the category where a single conversation, read on its own, tells you almost nothing useful.

AI-assisted moderation, in the way I actually used it, was never about handing a conversation over to a model and letting it decide anything. It was closer to feeding it everything a person could reasonably miss when they're moving fast: full timestamps, complete user history, every prior incident, the entire shape of a case rather than the one message someone happened to flag. Fed that way, it could filter out what didn't matter, the back-and-forth insults and noise that read as violations but weren't actually the substance of the case, and surface the pattern underneath faster than manually rereading everything myself would have.

That's a specific kind of help, and it only worked because of what I put into it, not because of what it figured out on its own. I gave it the full context, the history, the timestamps, along with a solid grounding in the platform's actual rules, the Discord TOS, and our own server guidelines. It wasn't operating on vague instinct about what might be a violation. It was working from the same complete picture I had, just faster at compressing it. Once I was fully across a case myself, I'd bring my own read and a proposed action, and we'd essentially work out the enforcement approach together from there, rather than me asking it what to do and taking the answer at face value.

That said, it wasn't flawless, and the failure mode that did show up was a specific one: occasionally mixing up user identities and histories, applying a standard penalty meant for one person's record to a case that actually belonged to someone else entirely. That's a real risk with this kind of tool, and it's exactly why the final read always stayed mine. I never treated its output as the decision. It was an input I checked against the actual case before anything went out, which is the only reason a mix-up like that never turned into an actual wrong call against the wrong person.

What actually changed because of this wasn't the standard I was enforcing to. It was how fast I could get from a messy,

drawn-out conflict to a defensible decision, especially in the cases most likely to turn into a fight, where speed and a documented trail mattered as much as getting the call right in the first place.

Suggested citation: Taghiev, T. "What AI-Assisted Moderation Actually Caught." Field Note, 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>