

Explainable Moderation: The Same Explanation Doesn't Work the Same Way Twice

Most writing about explainable moderation focuses on how much detail a rule violation notice should include. Less attention goes to tone, and to the fact that tone can't be standardized across communities without cost. Drawing on moderating the same game across Turkish, English, and Brazilian Portuguese communities, this piece treats tone as a resource allocation decision rather than a stylistic one.

Most writing about explainable moderation asks the same question: how much detail should a rule violation notice include. What gets less attention is that the tone of that explanation matters as much as its content, and that tone isn't something you can standardize across communities without cost.

I moderated Turkish, English, and Brazilian Portuguese communities for the same game, often issuing the same category of penalty across all three in a single shift. The written policy was identical. The delivery wasn't, and it wasn't supposed to be.

With Turkish users, I defaulted to neutral, low-affect language, deliberately. Almost any enforcement action, regardless of how clearly justified, would draw an appeal. Not a quiet one. Users would argue they'd been provoked, claim they didn't know the rule, or resurface the same complaint months later insisting the ban was unfair and they'd "done nothing." Some of that pushback stayed civil. Some of it escalated into insults, and a smaller number into outright fabrication, users telling other servers a moderator had a personal vendetta against them, reframing a justified penalty as targeted harassment. Warmth in that environment wasn't a neutral choice. It read as an opening. So the tone I used wasn't really about the individual case. It was risk management: a flat, procedural explanation gave less material to escalate, and conserved energy I'd need for the cases that actually required it.

With English-speaking users, the same neutral tone would have been a mismatch. That community ran calmer by default, disagreement without the personal-vendetta narrative, and enforcement was more often met with quiet acceptance or an actual apology. Being warmer here wasn't a risk. It was accurate to the environment, and it also did something specific: as my role became more visible, that same warmth read differently than it would have in the Turkish community. Authority read as approachable rather than as an opportunity to escalate.

The Brazilian Portuguese community is the clearest test of whether this was really about language or about something else, because I never spoke the language well enough to modulate tone through word choice alone. Every explanation went through translation. What I could control was intent behind the interaction, staying warm and operationally friendly regardless of the specific words landing correctly. I have not had a single user push back on an enforcement action in that community. That's not because the moderation was lighter. It's because the posture behind the explanation, even filtered through translation, seemed to matter more than its precision.

None of this means one tone is objectively correct and the others are overcautious. It means "explainable moderation" as a design principle assumes there's a single right way to explain a decision, when what I was actually managing across three communities was three different appeal cultures, three different baseline relationships to authority, and three different costs of getting the tone wrong. A flat explanation isn't inherently better or worse than a warm one. It's a resource allocation decision, made under the assumption that moderator attention and goodwill are finite, and that spending them uniformly across communities with radically different escalation patterns isn't neutral. It's just spending them in the wrong place.