


Governance Reporting from the Frontline: What a Report Looks Like When No One Acts On It

Most writing about governance reporting focuses on what a good report contains. Less gets said about what happens when a report is accurate, repeated, and correct, and still doesn't change anything. Drawing on four consecutive months of real moderation reports for a community I moderated without speaking its language, this piece looks at what reporting can actually force, and what it can only document.

A recommendation, repeated correctly with mounting evidence, is not the same thing as a decision — reporting can document a gap, but it can't force the fix.

Most writing about governance reporting focuses on what a good report contains. Less gets said about what happens when a report is good, repeated, and correct, and still doesn't change anything. I spent four consecutive months writing monthly moderation reports for a Brazilian Portuguese game community I moderated without speaking Portuguese. What that produced wasn't just a data problem. It was a real-time case study in what a report can and can't do on its own.

The pattern was visible from month one. Every report followed the same structure: activity data, violation breakdown, operational reality, recommendations. The recommendations were consistent across all four months, almost word for word: hire a Portuguese-speaking moderator, build a blacklist for patterns detectable without translation, establish a coordination protocol for high-activity periods. March called moderator hiring “not a comfort decision, a capacity decision.” June called it “not a comfort decision. It is a capacity decision” again, three months later, because nothing had changed.



Volume spiked in April, then eased back toward where it started (rounded)

The data behind the repetition tells its own story. Message volume roughly doubled between March and April (about 6,300 to about 11,900), then eased back through May (about 8,900) and June (about 5,700), nearly returning to where it started. Traffic came back down. The pattern underneath it didn't. Violations followed the same curve: April's

report logged 49 violations across categories from account trading to unauthorized recruitment; May counted 74, with account trading alone making up a third of all cases. By June, the violation types hadn't diversified so much as organized: account and gold trading, gambling-site referrals, and a pay-for-task scheme built around TikTok Lite referral links, structured enough to look like coordinated outside activity rather than isolated bad behavior.

What made repeat offenders hard to catch wasn't the language barrier alone. It was the absence of a system to remember them. The April report documents a user who committed the same violation six consecutive days, each time processed in isolation as a one-hour mute, because there was no tracking mechanism fast enough to catch the pattern under single-moderator conditions. Standard escalation logic (warning, then temporary mute, then longer mute, then ban) would have ended the behavior on day three. Instead it ran for six, not because the moderation policy was wrong, but because applying it correctly required remembering a case that no system was recording. May's report is explicit about this: nine separate ban evasion cases that month were caught only because one person recognized writing style and message patterns from memory. That's not a detection system. That's a moderator functioning as the detection system, and the report says so directly.

The reports also captured something reports don't usually have room for: who else was absorbing the gap. May's report includes a section crediting an unofficial volunteer who had been directing users to the right channels and explaining rules in two languages, alongside a direct statement that this was "a serious structural problem," not a success story. A community filling in for a missing moderation function isn't evidence the system is working. It's evidence of exactly how much slack the current setup depends on to look stable from the outside.

By June, the reports had started measuring their own failure to produce change. Each month's conclusion table scored the same categories: operational stability (maintained), violation detection (consistent), capacity sustainability (at critical limit), language barrier (not resolved), system design (incomplete). The status columns didn't move for four months. What did move was the language around them, from "close to breaking" in March to "past critical limit" in April. A report can hold a system's status constant in its own summary language while the substance underneath it keeps deteriorating; the report's job was never to fix that, only to make it visible in numbers regularly enough that "we didn't know" stopped being a defensible response.

What this points to. A recommendation, repeated correctly with mounting evidence, is not the same thing as a decision. Reporting can establish that a gap exists, quantify how it's growing, and document what specifically breaks when it's not closed, month after month, on a fixed schedule, in language leadership can act on. What it can't do is force the action itself. That's a different function, with a different owner, and conflating the two is exactly how an operational report turns into an accountability record instead: not proof that the system worked, but proof of exactly when someone knew it wasn't going

to.

Suggested citation: Taghiev, T. "Governance Reporting from the Frontline." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at <https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>