

Repeat Offenders: The Violation Is the Pattern, Not the Message

A month of hand-logged moderation data showed most violators appear once and disappear, while a small group returns again and again. Per-incident moderation handles every case correctly and still gets the person wrong.

The unit of moderation is the incident, but the unit of risk is the person — and anywhere that gap is bridged only by a moderator's memory, the community is one staffing change away from losing its pattern detection.

Most moderation writing treats a violation as an event: a message appears, it breaks a rule, someone acts on it, the case closes. The queue is built around that unit. So are the logs, the reports, and most of the tooling. After a month of moderating a Brazilian Portuguese game community and logging every violation by hand, I think that unit is wrong, or at least incomplete. The message is what you act on. The person is where the risk lives.

Here is what a month actually looked like. Across thirty days and about 5,700 messages, I logged roughly 130 violation cases. They came from about 70 unique users. That ratio already tells you something: most people who break a rule do it once, get warned or muted, and never show up in the log again. But roughly one in five violators came back at least once. And a much smaller group came back again and again, across different days, with the same behavior each time.

One case made the problem impossible to ignore. A user asked to buy an account in the channel. Account trading was a clear violation, so the message was removed and the user warned. The next day, the same request, twice. The day after, again. By the time the same user posted the same request on a fifth separate day, the individual messages had not changed at all: each one, taken alone, was a mild, almost polite violation, the kind that earns a routine warning. What had changed was everything around them. This was no longer a user who broke a rule. It was a user whose relationship to the rule was the violation.

Per-incident moderation handles each of those days correctly and still gets the month wrong. Every removal was justified. Every warning followed procedure. If you audited any single case, you would find nothing to criticize. But the correct handling of five incidents added up to an incorrect handling of one person, because the escalation logic lived in my memory, not in the system. I recognized the name. I remembered the previous days. The ban recommendation that eventually went forward was built on continuity that existed nowhere except in a moderator who happened to have been on shift all week.

That is the structural weakness. A moderation system that tracks incidents but not people delegates its pattern detection to whoever happens to be watching. It works when the same moderator covers the same channel every day, holds the names in their head, and writes careful logs. It fails silently the moment coverage rotates, volume spikes, or the moderator who carries the informal database leaves. The pattern does not disappear. The ability to see it does.

The repeat cases also behaved differently from the one-off cases, in a way that matters for enforcement design. First offenses in that month were scattered across categories: spam, self-promotion, off-topic begging, an occasional heated exchange. The repeat offenses concentrated. They clustered heavily around account trading and adjacent commerce: buying, selling, matching buyers with sellers. At one point a second user started tagging the repeat offender to connect a purchase. That is not two isolated violations. That is a small market forming, and a market is a pattern by definition. You cannot see it at the message level, because each message is just one side of one conversation.

This changed how I wrote recommendations. A first-time violator got the action the incident deserved. A repeat name got the action the trajectory deserved, and I started saying so explicitly in the log: not "user posted a prohibited request" but "user posted the same prohibited request for the fifth day; warnings have demonstrably not changed the behavior; recommend removal." The evidence for a ban was never the fifth message. It was the failure of four previous interventions, documented in sequence.

There is an obvious objection: giving weight to history risks punishing people for who they are rather than what they did. It is a fair concern, and it is why the sequence matters. Escalation earned by documented, repeated, same-category violations is not profiling. It is the opposite: it is the system giving someone four chances and recording each one. What would actually be unfair is the alternative that per-incident moderation quietly produces, where a persistent bad actor and a one-time offender receive identical treatment forever because the system cannot tell them apart.

What this points to. The unit of moderation is the incident, but the unit of risk is the person, and most community tooling only knows about the first. Anywhere the gap between those two is bridged by an individual moderator's memory, the community is one staffing change away from losing its pattern detection entirely. If I could push one requirement into any moderation stack, it would not be a better filter. It would be a persistent, queryable answer to the oldest question in the queue: have we seen this person before, and for what?

Suggested citation: Taghiev, T. "Repeat Offenders: The Violation Is the Pattern, Not the Message." Field Note, July 2026.

Turan Taghiev writes on trust & safety, platform governance, and AI governance, drawing on direct moderation and technical-support experience across multiple language communities. More at

<https://deheidhemklashwo-dev.github.io/TuranTaghiev.github.io/>